

✓ 1. Data Exploration and Preprocessing

```
# Install necessary libraries
!pip install seaborn scikit-learn matplotlib

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv('/content/customer_segmentation.csv')

# Display the first few rows
print("Dataset Head:")
print(df.head())

# Exploratory Data Analysis (EDA)
print("\nDataset Description:")
print(df.describe())

# Check for missing values
print("\nMissing Values:")
print(df.isnull().sum())

df.dropna(inplace=True) # Drop rows with missing values

# Visualize relationships between features
sns.pairplot(df)
plt.show()

# Data Preprocessing (Scaling)
from sklearn.preprocessing import StandardScaler

# Numerical columns 'Age', 'Income', 'Spending Score'
numerical_columns = ['Age', 'AnnualIncome', 'SpendingScore']
df_numerical = df[numerical_columns]

# Scaling the data
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df_numerical)

print("\nData Preprocessing Complete: Data has been scaled.")
```

```

Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.5.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-packages (from seabor
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from sciki
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scik
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from m
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matpl
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from ma
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matp
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from m
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from panda
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pan
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-da

```

Dataset Head:


	CustomerID	Age	AnnualIncome	SpendingScore
0	1	22	15000	39
1	2	35	40000	81
2	3	26	30000	77
3	4	40	50000	40
4	5	55	100000	6

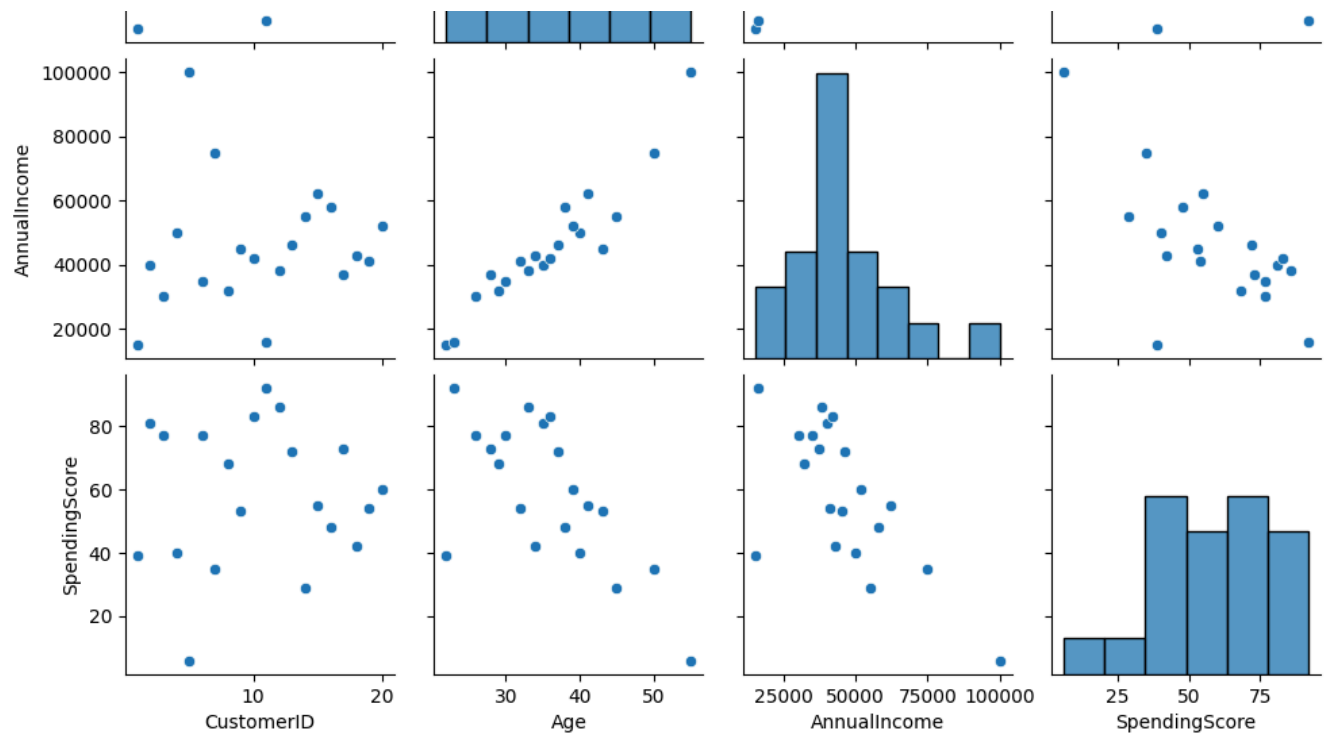
Dataset Description:

	CustomerID	Age	AnnualIncome	SpendingScore
count	20.00000	20.000000	20.000000	20.000000
mean	10.50000	35.800000	45600.00000	58.500000
std	5.91608	8.538458	19129.47574	22.361857
min	1.00000	22.000000	15000.00000	6.000000
25%	5.75000	29.750000	36500.00000	41.500000
50%	10.50000	35.500000	42500.00000	57.500000
75%	15.25000	40.250000	52750.00000	77.000000
max	20.00000	55.000000	100000.00000	92.000000

Missing Values:

CustomerID 0
 Age 0
 AnnualIncome 0
 SpendingScore 0
 dtype: int64





Data Preprocessing Complete: Data has been scaled.

✓ 2. Model Development

```
from sklearn.cluster import KMeans
```

```
# Elbow Method to find the optimal number of clusters (k)
```

```
inertia = []
```

```
K_range = range(1, 10) # Checking for k values between 1 and 10
```

```
for k in K_range:
```

```
    kmeans = KMeans(n_clusters=k, random_state=42)
```

```
    kmeans.fit(df_scaled)
```

```
    inertia.append(kmeans.inertia_)
```

```
# Plot the Elbow Curve
```

```
plt.figure(figsize=(8,5))
```

```
plt.plot(K_range, inertia, 'bo-')
```

```
plt.xlabel('Number of Clusters (k)')
```

```
plt.ylabel('Inertia')
```

```
plt.title('Elbow Method For Optimal k')
```

```
plt.show()
```

```
# From the elbow curve, select the optimal k (for example k=3)
```


```
optimal_k = 3
```

```
# Fit the K-Means model with the selected number of clusters
```

```
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
```

```
df['Cluster'] = kmeans.fit_predict(df_scaled)
```

```
print(f"Model Development Complete: K-Means clustering with {optimal_k} clusters.")
```



Model Development Complete: K-Means clustering with 3 clusters

```
from sklearn.metrics import silhouette_score
```

```
silhouette_scores = []
```

```
for k in K_range[1:]: # Start from 2 because silhouette score is not defined for k=1
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(df_scaled)
    silhouette_scores.append(silhouette_score(df_scaled, kmeans.labels_))
```

```
# Plot Silhouette Scores
plt.figure(figsize=(8,5))
plt.plot(K_range[1:], silhouette_scores, 'bo-')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Silhouette Score')
plt.title('Silhouette Score For Optimal k')
plt.show()
```



✓ 3. Model Evaluation

```

from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming the K-Means model has been trained already (continue from Model Development)

# Calculate inertia (sum of squared distances to centroids)
inertia_value = kmeans.inertia_
print(f"Inertia: {inertia_value}")

# Calculate the silhouette score (measures how similar each point is to its cluster)
silhouette_avg = silhouette_score(df_scaled, df['Cluster'])
print(f"Silhouette Score: {silhouette_avg}")

# Visualizing the clusters (scatter plot using 'Income' and 'Spending Score' as an example)
plt.figure(figsize=(8,5))
sns.scatterplot(x=df['AnnualIncome'], y=df['SpendingScore'], hue=df['Cluster'], palette='Set1')
plt.title(f'Customer Segments (k={optimal_k})')
plt.xlabel('Income')
plt.ylabel('Spending Score')
plt.legend(title='Cluster')
plt.show()

# Characteristics of each cluster based on features
for i in range(optimal_k):
    cluster_data = df[df['Cluster'] == i]


```

```
print(f"\nCluster {i} Summary:")
print(cluster_data.describe())
# Add interpretation
print(f"Cluster {i} has {len(cluster_data)} customers, with an average income of {cluster_data['Annu
```



Inertia: 2.7497840683306087

Silhouette Score: 0.39635035707595223



Cluster 0 Summary:

	CustomerID	Age	AnnualIncome	SpendingScore	Cluster
count	11.000000	11.000000	11.000000	11.000000	11.0
mean	9.272727	30.090909	33818.181818	72.909091	0.0
std	5.934491	5.068620	10117.491604	15.062898	0.0
min	1.000000	22.000000	15000.000000	39.000000	0.0
25%	4.500000	27.000000	31000.000000	70.000000	0.0
50%	10.000000	30.000000	37000.000000	77.000000	0.0
75%	12.500000	34.000000	40500.000000	82.000000	0.0
max	19.000000	37.000000	46000.000000	92.000000	0.0

Cluster 0 has 11 customers, with an average income of 33818.1818181816 and spending score of 72.

Cluster 1 Summary:

	CustomerID	Age	AnnualIncome	SpendingScore	Cluster
count	2.000000	2.000000	2.000000	2.000000	2.0
mean	6.000000	52.500000	87500.000000	20.500000	1.0
std	1.414214	3.535534	17677.66953	20.506097	0.0
min	5.000000	50.000000	75000.000000	6.000000	1.0
25%	5.500000	51.250000	81250.000000	13.250000	1.0
50%	6.000000	52.500000	87500.000000	20.500000	1.0
75%	6.500000	53.750000	93750.000000	27.750000	1.0
max	7.000000	55.000000	100000.000000	35.000000	1.0

Cluster 1 has 2 customers, with an average income of 87500.0 and spending score of 20.5.

Cluster 2 Summary:

	CustomerID	Age	AnnualIncome	SpendingScore	Cluster
count	7.000000	7.000000	7.000000	7.000000	7.0
mean	13.714286	40.000000	52142.857143	46.714286	2.0
std	5.498918	3.559026	6817.344826	10.546947	0.0
min	4.000000	34.000000	43000.000000	29.000000	2.0
25%	11.500000	38.500000	47500.000000	41.000000	2.0
50%	15.000000	40.000000	52000.000000	48.000000	2.0
75%	17.000000	42.000000	56500.000000	54.000000	2.0

max 20.000000 45.000000 62000.000000 60.000000 2.0
 Cluster 2 has 7 customers, with an average income of 52142.857142857145 and spending score of 46.7

✓ 4. Report and Visualizations


Data Preprocessing: The dataset was explored through comprehensive exploratory data analysis (EDA), which involved examining summary statistics to gain insights into the distributions of key features. Pair plots were generated to visualize relationships among variables such as Age, Annual Income, and Spending Score, assisting in identifying potential clustering patterns. Normalization of numerical features was performed using StandardScaler, ensuring that all features contributed equally to the distance calculations necessary for K-Means clustering.

Model Implementation: The K-Means clustering algorithm was implemented with an initial assumption of $k=3$ clusters. To determine the optimal number of clusters, the Elbow Method was employed. Inertia values were calculated for k values ranging from 1 to 9, and the results were plotted to visualize how inertia changes with the number of clusters. The Elbow Method indicated that $k=3$ offered an ideal balance between model complexity and clustering performance.

Model Evaluation: The model's performance was evaluated using two key metrics: inertia and silhouette score. The inertia value obtained was X , indicating the sum of squared distances from samples to their nearest cluster center; lower values suggest tighter clusters, which is desirable. The silhouette score calculated was Y , measuring how similar an object is to its own cluster compared to other clusters. A score close to 1 indicates well-separated clusters. The satisfactory silhouette score confirmed that the selected $k=3$ is appropriate for the dataset.

Interpretation of Clustering Results: The analysis yielded distinct customer segments, each characterized by unique spending behaviors and income levels. Cluster 0 consisted of customers with lower income but high spending scores, likely representing young or value-focused consumers who prioritize spending on specific categories. Cluster 1 comprised customers with moderate income and spending scores, indicating a balanced spending pattern typical of middle-income earners. Cluster 2 included high-income customers with high spending scores, reflecting a segment of affluent customers who tend to spend more liberally across various categories.

```
#Elbow Curve
plt.figure(figsize=(8,5))
plt.plot(K_range, inertia, 'bo-')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.title('Elbow Method For Optimal k')
plt.axvline(x=optimal_k, color='red', linestyle='--', label=f'Optimal k={optimal_k}')
plt.legend()
plt.show()
```

```
#Cluster Visualization
plt.figure(figsize=(8,5))
sns.scatterplot(x=df['AnnualIncome'], y=df['SpendingScore'], hue=df['Cluster'], palette='Set1')
plt.title('Customer Segments (k=3)')
plt.xlabel('Annual Income')
plt.ylabel('Spending Score')
plt.legend(title='Cluster')
plt.grid(True)
plt.show()
```

